

REPORT.PROJECT

Nguyen Ngoc Son 009, M2-ICT USTH

08/01/2023

Summary

In this lab work, CPU and 2D GPU parallelism are implemented with Kuwahara filter. The framework used is numba in python under Google Colab GPU session (Tesla GPU).



Figure 1: Original RGB photo

As you can see in the figure 1, the image is RGB originally. After implementation of Kuwahara code with CPU/GPU, the resulting image is an animated picture.

Implementations

Listing 1: Sample Python code – gray scale RGB image 1D and 2D implementations.

```
1
2 # CPU implementation
3 def Kuwahara(image, size, v_ = None):
4
```

```

5     image = image.astype(np.float64)
6     output = np.zeros(image.shape)
7     mean = np.zeros([4, image.shape[0], image.shape[1]])
8     std = mean.copy()
9
10    kernel_template = np.hstack((np.ones((1, int((size-1)/2)+1)), np.↵
        zeros((1, int((size-1)/2)))))
11    # =====
12    # [[1. 1. 1. 0. 0.]]
13    pad = np.zeros((1, size))
14    # =====
15    # [[0. 0. 0. 0. 0.]]
16    kernel = np.tile(kernel_template, (int((size-1)/2)+1, 1))
17    # =====
18    # [[1. 1. 1. 0. 0.]
19    # [1. 1. 1. 0. 0.]
20    # [1. 1. 1. 0. 0.]]
21    kernel = np.vstack((kernel, np.tile(pad, (int((size-1)/2), 1))))
22    # =====
23    # [[1. 1. 1. 0. 0.]
24    # [1. 1. 1. 0. 0.]
25    # [1. 1. 1. 0. 0.]
26    # [0. 0. 0. 0. 0.]
27    # [0. 0. 0. 0. 0.]]
28    # =====
29    average = kernel/np.sum(kernel)
30    # =====
31    # [[0.11111111 0.11111111 0.11111111 0.          0.          ]
32    # [0.11111111 0.11111111 0.11111111 0.          0.          ]
33    # [0.11111111 0.11111111 0.11111111 0.          0.          ]
34    # [0.          0.          0.          0.          0.          ]
35    # [0.          0.          0.          0.          0.          ]]
36    # =====
37
38    kernelstack = np.empty((4, size, size))
39    kernelstack[0] = average                                # a
40    kernelstack[1] = np.fliplr(average)                    # b
41    kernelstack[2] = np.flipud(average)                    # c
42    kernelstack[3] = np.fliplr(kernelstack[2])            # d
43
44    # [ a  a  ab  b  b]
45    # [ a  a  ab  b  b]
46    # [ac ac abcd bd bd]
47    # [ c  c  cd  d  d]
48    # [ c  c  cd  d  d]
49
50    for i in range(4):

```

```

51         mean[i] = conv2d(image, kernelstack[i]) #↵
           mean
52         std[i] = conv2d(image**2, kernelstack[i]) - mean[i]**2 #↵
           variance
53
54     if v_ is not None:
55         indices = np.argmin(v_,0)
56         for i in range(image.shape[0]):
57             for k in range(image.shape[1]):
58                 output[i,k] = mean[indices[i,k], i,k].astype('uint8')
59     else:
60         indices = np.argmin(std,0)
61
62         for i in range(image.shape[0]):
63             for k in range(image.shape[1]):
64                 output[i,k] = mean[indices[i,k], i,k]
65
66     return output, std

```

Listing 2: Averaging Kernel preparation

```

1
2     # =====
3     # [[0.11111111 0.11111111 0.11111111 0.          0.          ]
4     # [0.11111111 0.11111111 0.11111111 0.          0.          ]
5     # [0.11111111 0.11111111 0.11111111 0.          0.          ]
6     # [0.          0.          0.          0.          0.          ]
7     # [0.          0.          0.          0.          0.          ]]
8     # =====
9
10    kernelstack = np.empty((4,size,size))
11    kernelstack[0] = average # a
12    kernelstack[1] = np.fliplr(average) # b
13    kernelstack[2] = np.flipud(average) # c
14    kernelstack[3] = np.fliplr(kernelstack[2]) # d
15
16    # [ a  a  ab  b  b]
17    # [ a  a  ab  b  b]
18    # [ac ac abcd bd bd]
19    # [ c  c  cd  d  d]
20    # [ c  c  cd  d  d]

```

Averaging kernel:

In this method we do not need to run pixel by pixel but rather convolve the image!

In turn, all of the filters (a, b, c, d) will be prepared as above and multiply with the

padded images, resulting in std calculation, RGB output.

To be able to get STD of V in HSV we implement the RGB to HSV conversion code:

Listing 3: rgb to hsv

```
1
2 def rgb_to_hsv(r, g, b):
3     r, g, b = r/255.0, g/255.0, b/255.0
4     mx = max(r, g, b)
5     mn = min(r, g, b)
6     df = mx-mn
7     # print(mx)
8     if mx == mn:
9         h = 0.0
10    elif mx == r:
11        h = 60*((g-b)/df % 6)
12    elif mx == g:
13        h = 60*((b-r)/df + 2)
14    elif mx == b:
15        h = 60*((r-g)/df + 4)
16    if mx == 0.0:
17        s = 0.0
18    else:
19        s = (df/mx)
20    v = mx
21    return h, s, v
```

GPU implementation code

Listing 4: RGB to HSV conversion with GPU

```
1
2 @cuda.jit
3 def rgb_to_hsv(in_, out):
4     x = cuda.threadIdx.x + cuda.blockIdx.x * cuda.blockDim.x
5     y = cuda.threadIdx.y + cuda.blockIdx.y * cuda.blockDim.y
6     r, g, b = in_[x, y, 0], in_[x, y, 1], in_[x, y, 2]
7     mx = max(r, g, b)
8     mn = min(r, g, b)
9     df = mx-mn
10
11    if mx == mn:
12        h = 0.0
13    elif mx == r:
```



Figure 2: On the left: Original image, On the right: Filtered image

```

14         h = 60*((g-b)/df % 6)
15
16     elif mx == g:
17         h = 60*((b-r)/df + 2)
18
19     elif mx == b:
20         h = 60*((r-g)/df + 4)
21
22     if mx == 0.0:
23         s = 0.0
24     else:
25         s = (df/mx)
26     v = mx
27
28     out[0, x, y] = h
29     out[1, x, y] = s
30     out[2, x, y] = v

```

RGB to HSV in GPU

Here we keep the similar code but manage the memory of GPU to run in parallel.

Listing 5: Kuwahara using V value from HSV

```

2  @cuda.jit
3  def Kuwahara_v(RGB, V_in, KuwaRGB, kernelSize):
4      tidx = cuda.threadIdx.x + cuda.blockIdx.x * cuda.blockDim.x
5      tidy = cuda.threadIdx.y + cuda.blockIdx.y * cuda.blockDim.y
6      # Memory management is the most counter-intuitive part of this ↵
7      project
8      kernel_map = (
9          ((tidx-kernelSize, tidx+1), (tidy-kernelSize, tidy+1))↵
10         ,
11         ((tidx, tidx+kernelSize+1), (tidy-kernelSize, tidy+1))↵
12         ,
13         ((tidx-kernelSize, tidx+1), (tidy, tidy+kernelSize+1))↵
14         ,
15         ((tidx, tidx+kernelSize+1), (tidy, tidy+kernelSize+1))
16     )
17     min_std = 99999.0
18     min_idx = 0
19     for idx in range(4):
20         # Previously this part was CPU implementation.
21         # mean[i] = conv2d(image, kernelstack[i]) #↵
22         mean
23         # std[i] = conv2d(image**2, kernelstack[i]) - mean[i]**2 #↵
24         variance
25         # y, x = image.shape
26         # y = y - height + 1
27         # x = x - height + 1
28         # new_image = np.zeros((y ,x))
29         # print (new_image.shape)
30
31         # Now convolution and STD calculation will be conducted in ↵
32         parallel as below
33
34         sum = 0.0
35         sumSquare = 0.0
36         for wi in range(*kernel_map[idx][0]):
37             for wj in range(*kernel_map[idx][1]):
38                 sum += V_in[2,wi, wj]
39                 sumSquare += V_in[2, wi, wj] **2
40
41         mean = sum/(kernelSize+1)**2
42         std = math.sqrt(abs(sumSquare /(kernelSize+1)**2 - mean**2))
43         if std < min_std:
44             min_std = std
45             min_idx = idx
46
47     # After STD calculation we can use the V with minimum STD to get ↵
48     the average of RGB value inside of the kernel

```

```

41     sum_r = 0.0
42     sum_g = 0.0
43     sum_b = 0.0
44     for i in range(*kernel_map[min_idx][0]):
45         for j in range(*kernel_map[min_idx][1]):
46             sum_r += RGB[i, j, 0]
47             sum_g += RGB[i, j, 1]
48             sum_b += RGB[i, j, 2]
49
50     # Average of RGB with V index and return to output
51     KuwaRGB[tidx, tidy, 0] = sum_r / (kernelSize+1)**2
52     KuwaRGB[tidx, tidy, 1] = sum_g / (kernelSize+1)**2
53     KuwaRGB[tidx, tidy, 2] = sum_b / (kernelSize+1)**2

```

Kuwahara in GPU

Here the result ! a little too big kernel size! Memory management is the most counter-intuitive part of this project, while convolution in GPU is the most mind bending experience!



Figure 3: On the left: Original image, On the right: Filtered image with GPU